


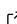


# simChef: High-quality data science simulations in R

James Duncan <sup>1\*</sup>, Tiffany Tang <sup>2\*</sup>, Corrine F. Elliott <sup>2</sup>, Philippe Boileau <sup>1</sup>, and Bin Yu<sup>1,2</sup>

<sup>1</sup> Graduate Group in Biostatistics, University of California, Berkeley <sup>2</sup> Department of Statistics, University of California, Berkeley  Corresponding author \* These authors contributed equally.

DOI: [N/A](#)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: 01 January 1970

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

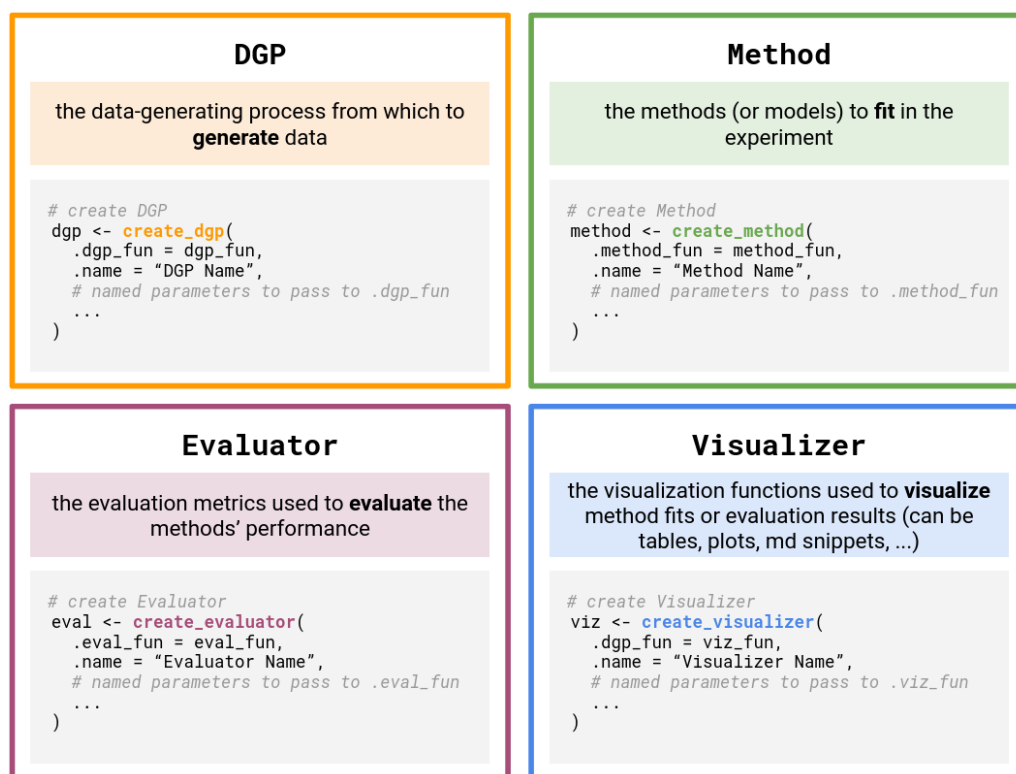
## Summary



Data science simulation studies occupy an important role in data science research as a means to gain insight into new and existing statistical methods. Whether as a means to establish comprehensive benchmarks of existing procedures for a common task, to demonstrate the strengths and weaknesses of novel methodology applied to synthetic and real-world data, or to probe the validity of a theoretical analysis, simulations serve as statistical sandboxes that open a path toward otherwise inaccessible discoveries. Yet creating high-quality simulation studies typically involves a number of repetitive and error-prone coding tasks, such as implementing data-generating processes (DGPs) and statistical methods, sampling from these DGPs, parallelizing computation of simulation replicates, summarizing metrics, and visualizing, documenting, and saving results. While this administrative overhead is necessary to reach the end goals of a given data science simulation, it is not sufficient, as the data scientist must navigate a number of important judgment calls such as the choice of data settings, baseline statistical methods, associated parameters, and evaluation metrics for scientific relevancy. The scientific context varies drastically from one study to the next while the simulation scaffolding remains largely similar; yet simulation code repositories often lack the flexibility to easily allow for reuse in novel settings or even simple extension when new questions arise in the original context.

simChef addresses the need for an intuitive, extensible, and reusable framework for data science simulations. Drawing substantially from the Predictability, Computability, and Stability (PCS) framework (Yu & Kumbier, 2020), simChef empowers data scientists to focus their attention toward the scientific best practices encompassed by PCS by removing many of the administrative burdens of simulation design with an intuitive [tidy grammar](#) of data science simulations and automated interactive R Markdown documentation.

## A powerful grammar of data science simulations



**Figure 1:** simChef provides four classes which implement distinct simulation objects in an intuitive and modular manner: DGP, Method, Evaluator, and Visualizer.

Inspired by the tidyverse (Wickham et al., 2019), simChef develops an intuitive grammar of simulation studies:

```
library(simChef)

dgp1 <- create_dgp(dgp_fun1, "my_dgp1", sd = 0.5)
dgp2 <- create_dgp(dgp_fun2, "my_dgp2")
method <- create_method(method_fun, "my_method")
eval <- create_evaluator(eval_fun)
viz <- create_visualizer(viz_fun)

exper <- create_experiment(dgp_list = list(dgp1, dgp2)) %>%
  add_method(method) %>%
  add_vary_across(
    list(dgp1, dgp2),
    n = c(1e2, 1e3, 1e4)
  ) %>%
  add_vary_across(
    dgp2,
    sparse = c(FALSE, TRUE)
  ) %>%
  add_vary_across(
    method,
    scalar_valued_param = c(0.1, 1.0, 10.0),
```

```
vector_valued_param = list(c(1, 2, 3), c(4, 5, 6)),
list_valued_param = list(list(a1=1, a2=2, a3=3),
                           list(b1=3, b2=2, b3=1))
) %>%
add_evaluator(eval) %>%
add_viz(viz)

future::plan(multicore, workers = 64)

results <- exper %>%
  run_experiment(n_reps = 100, save = TRUE)

new_method <- create_method(new_method_fun, 'my_new_method')

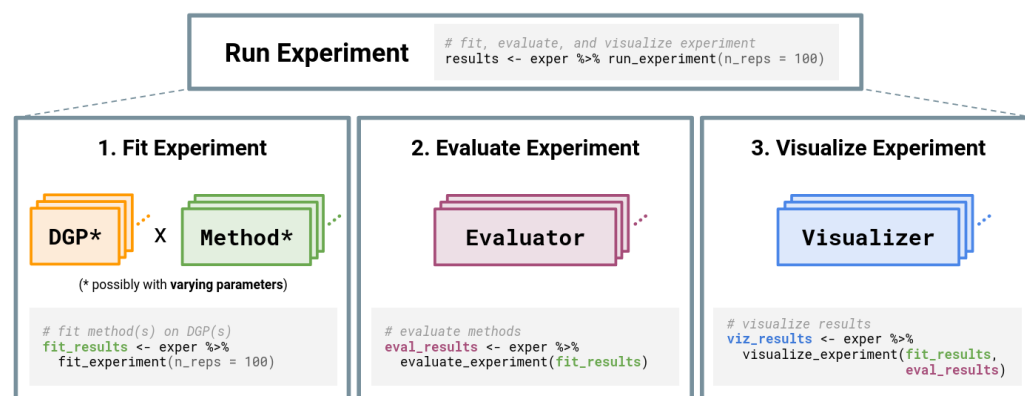
exper <- exper %>%
  add_method(new_method)

results <- exper %>%
  run_experiment(n_reps = 100, use_cached = TRUE)

init_docs(exper)
render_docs(exper)
```

Internally, `simChef` provides a modular conceptualization of data science simulations using four R6 ([Chang, 2022](#)) classes, portrayed [Figure 1](#): DGP, Method, Evaluator, and Visualizer. Users create or reuse custom functions (`dgp_fun`, `method_fun`, `eval_fun`, and `viz_fun` above) aligned with their scientific goals. The custom functions are then optionally parameterized and encapsulated in one of the corresponding classes via a `create_*` method together with optional constant parameters (e.g., `sd` above).

A fifth R6 class, `Experiment`, serves as a concrete implementation of the user's intent to answer a specific scientific question. The `Experiment` stores references to the first four objects along with the DGP and Method parameters that should be varied and combined during the simulation run. Parameters that are common across the users functions can be added jointly (as is the case for the `n` parameter to `dgp_fun1` and `dgp_fun2` above) and can have arbitrary data type (such as `scalar_valued_param` and `vector_valued_param` to `method_fun`).



**Figure 2:** The `Experiment` class handles relationships between the four classes portrayed in [Figure 1](#). Experiments may have multiple DGP and Method objects, which are combined across the Cartesian product of their varying parameters (represented by `\*`). Once computed, each `Evaluator` and `Visualizer` take in simulation replicates, while `Visualizer` additionally receives evaluation summaries.

The Experiment class flexibly handles the computation of simulation replicates in parallel using future (Bengtsson, 2021). The number of replicates per combination of DGP, Method, and parameters specified via `add_vary_across` is determined by the `n_reps` argument to `run_experiment` (Figure 2). Because replication happens at the per-combination level, the effective total number of replicates in the Experiment depends on the number of DGPs, methods, and varied parameters.

In the first call to `run_experiment` in the above example, there are two DGP instances, both of which are varied across three values of `n` and one of which is additionally varied across two values of `sparse`. This effectively results in nine distinct configurations for data generation. For the single Method in the experiment, we use three values of `scalar_valued_param`, two of `vector_valued_param`, and another two of `list_valued_param`, giving 12 distinct configurations. Thus, there are a total of 108 DGP-method-parameter combinations in the experiment, each of which is replicated 100 times. Figure 3 provides a detailed schematic of the `run_experiment` workflow, along with the expected inputs to and outputs from user-defined functions.

Users can also choose to save the experiment's results to disk by passing `save = TRUE` to `run_experiment`. Once saved, the user can add new DGP and Method objects to the experiment and compute additional replicates without re-computing existing results via the `use_cached` option. Considering the example above, when we add `new_method` and call `run_experiment` with `use_cached = TRUE`, `simChef` finds that the cached results are missing combinations of `new_method`, existing DGPs, and their associated parameters, giving nine new configurations. Replicates for the new combinations are then appended to the cached results.

Automated documentation in an interactive R Markdown template gathers the scientific details, summary tables, and visualizations side-by-side with the user's custom source code and parameters for data-generating processes, statistical methods, evaluation metrics, and plots. A call to `init_docs` generates empty markdown files for the user to populate with their overarching simulation objectives and with descriptions of each of the DGP, Method, Evaluator, and Visualizer objects included in the Experiment. Finally, a call to `render_docs` prepares the interactive R Markdown document, either for iterative design and analysis of the simulation or to provide a high-quality overview that can be easily shared. We provide an example of the simulation documentation at [this link](#) and corresponding source code is available on GitHub at [PhilBoileau/simChef-case-study](#).

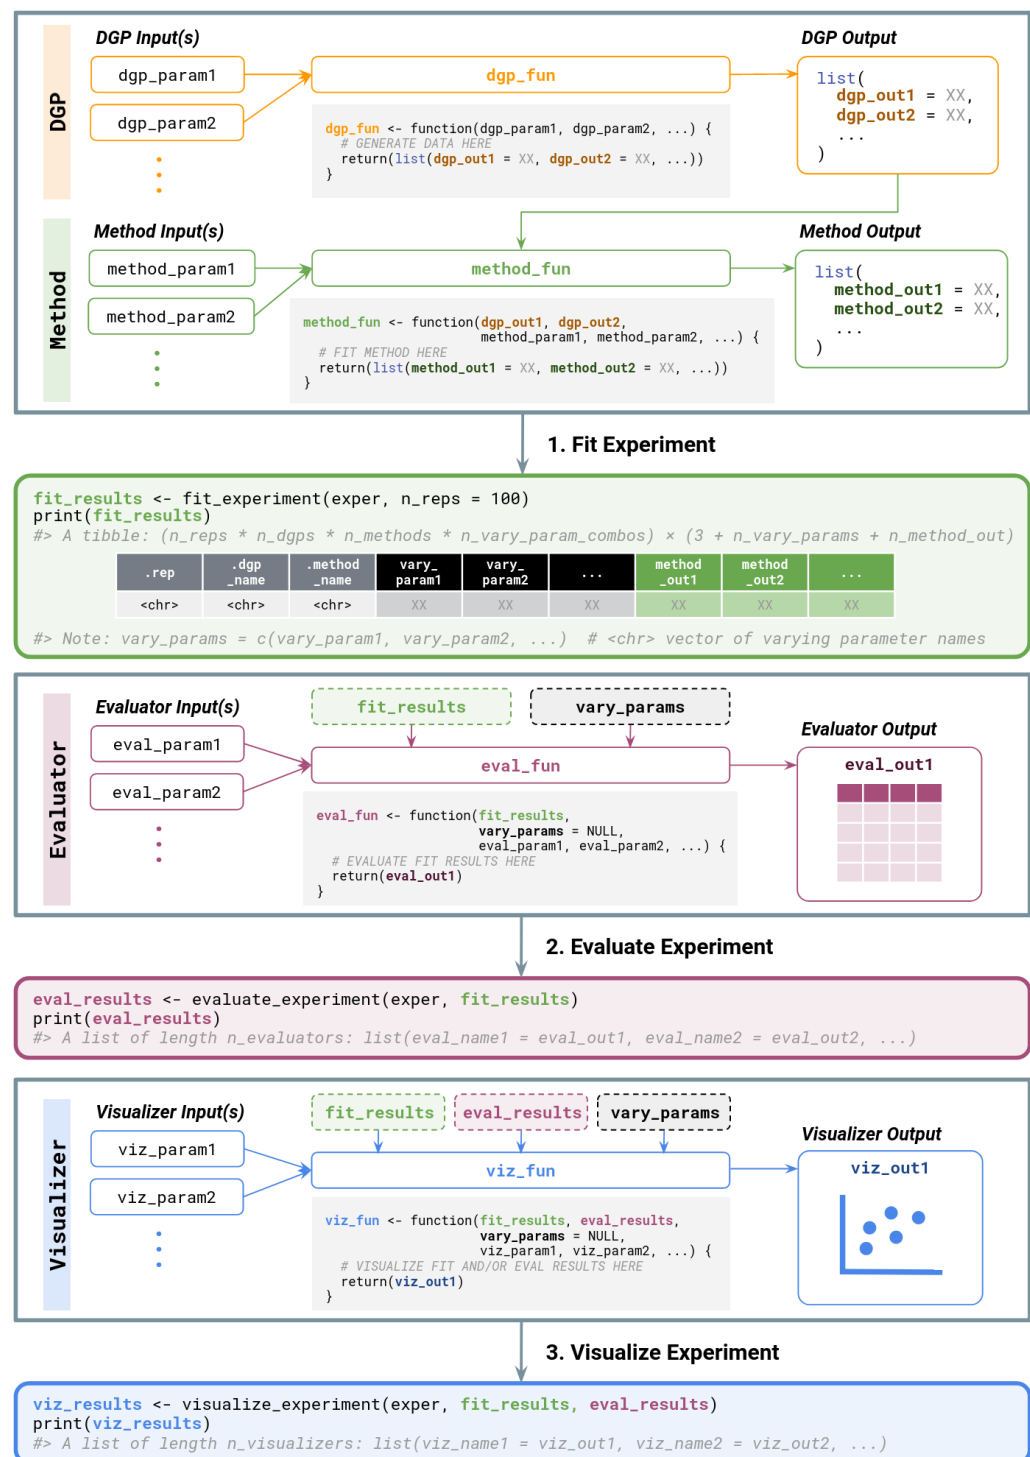


Figure 3: Detailed schematic of the run\_experiment workflow.

## Acknowledgements

## References

Bengtsson, H. (2021). A Unifying Framework for Parallel and Distributed Processing in R

using Futures. *The R Journal*, 13(2), 208. <https://doi.org/10.32614/RJ-2021-048>

Chang, W. (2022). *R6: Encapsulated classes with reference semantics*.

Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Grolemond, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M., Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., ... Yutani, H. (2019). Welcome to the Tidyverse. *Journal of Open Source Software*, 4(43), 1686. <https://doi.org/10.21105/joss.01686>

Yu, B., & Kumbier, K. (2020). Veridical data science. *Proceedings of the National Academy of Sciences*, 117(8), 3920–3929. <https://doi.org/10.1073/pnas.1901326117>