

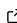


simChef: High-quality data science simulations in R

James Duncan ^{1*}, Tiffany Tang ^{2*}, Corrine F. Elliott ², Philippe Boileau ¹, and Bin Yu^{1,2,3,4}

¹ Graduate Group in Biostatistics, University of California, Berkeley ² Department of Statistics, University of California, Berkeley ³ Department of Electrical Engineering and Computer Sciences, University of California, Berkeley ⁴ Center for Computational Biology, University of California, Berkeley
¶ Corresponding author * These authors contributed equally.

DOI: [N/A](#)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: 01 January 1970

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary



Data science simulation studies occupy an important role in data science research as a means to gain insight into new and existing statistical methods. In particular, simulations serve as statistical sandboxes that open a path toward otherwise inaccessible discoveries. For example, they can be used to establish comprehensive benchmarks of existing procedures for a common task, to demonstrate the strengths and weaknesses of novel methodology applied to synthetic and real-world data, or to probe the validity of a theoretical analysis. Yet creating high-quality simulation studies typically involves a number of repetitive and error-prone coding tasks, such as implementing data-generating processes (DGPs) and statistical methods, sampling from these DGPs, parallelizing computation of simulation replicates, summarizing metrics, and visualizing, documenting, presenting, and saving results. While this administrative overhead is necessary to reach the end goals of a given data science simulation, it is not sufficient, as the data scientist must navigate a number of important judgment calls such as the choice of DGPs, baseline statistical methods, associated parameters, and evaluation metrics for scientific relevancy. The scientific context varies drastically from one study to the next while the simulation scaffolding remains largely similar; yet simulation code repositories often lack the flexibility to allow for facile reuse in novel settings or even for simple extension when new questions arise in the original context.

simChef addresses the need for an intuitive, extensible, and reusable framework for data science simulations. Drawing substantially from the Predictability, Computability, and Stability (PCS) framework (Yu & Kumbier, 2020), simChef empowers data scientists to focus their attention toward the scientific best practices encompassed by PCS by removing many of the administrative burdens of simulation design with an intuitive [tidy grammar](#) of data science simulations and automated interactive R Markdown documentation.

Core abstractions of data science simulations

At its core, `simChef` breaks down a simulation experiment into four modular components (Figure 1), each implemented as an R6 class (Chang, 2022):

- DGP: the data-generating processes from which to *generate* data
- Method: the methods (or models) to *fit* in the experiment
- Evaluator: the evaluation metrics used to *evaluate* the methods' performance
- Visualizer: the visualization functions used to *visualize* outputs from the method fits or evaluation results (can be tables, plots, or even R Markdown snippets to display)

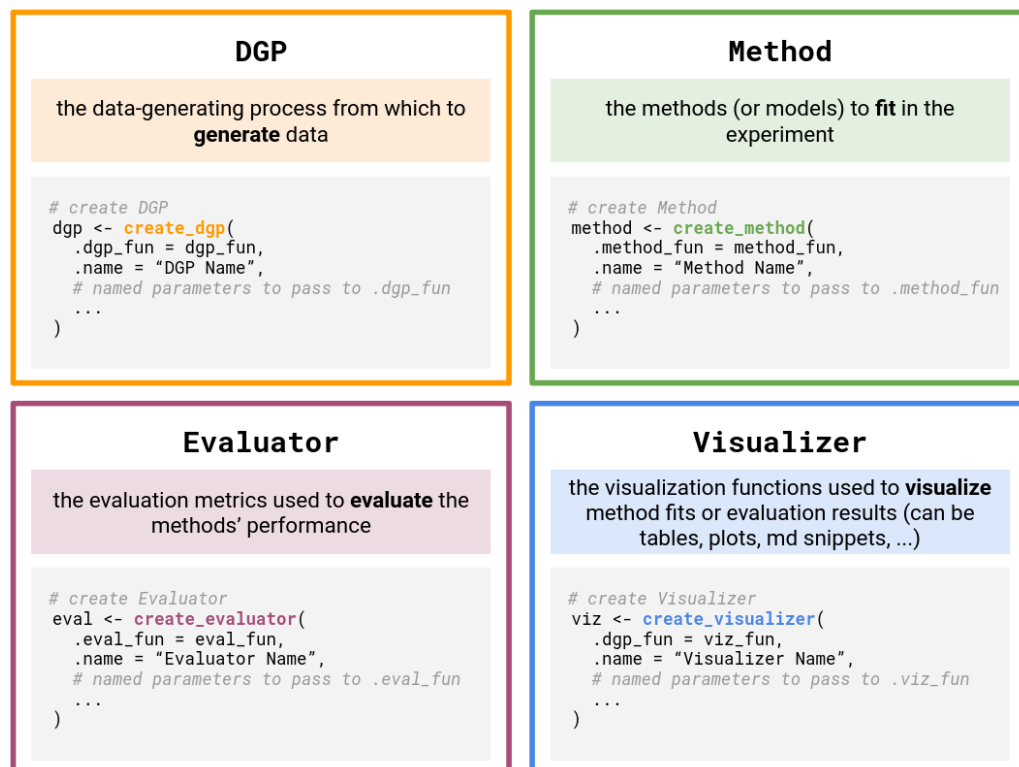


Figure 1: Overview of the four core components in a `simChef` Experiment. `simChef` provides four classes that implement distinct simulation objects in an intuitive and modular manner: DGP, Method, Evaluator, and Visualizer. Using these classes, users can easily build a `simChef` Experiment using reusable, customizable functions (i.e., `dgp_fun`, `method_fun`, `eval_fun`, and `viz_fun`). Optional named parameters can be set in these custom functions via the `...` arguments in the `create_*` methods.

Using these classes, users can create or reuse custom functions (i.e., `dgp_fun`, `method_fun`, `eval_fun`, and `viz_fun` in Figure 1) aligned with their scientific goals. The custom functions then can be parameterized and encapsulated in one of the corresponding classes via a `create_*` method, together with optional named parameters (see Figure 1).

A fifth R6 class, `Experiment`, unites the four components above and serves as a concrete implementation of the user's intent to answer a specific scientific question. Specifically, the `Experiment` stores references to the DGP(s), Method(s), Evaluator(s), and Visualizer(s) along with the DGP and Method parameters that should be varied and combined during the simulation run.

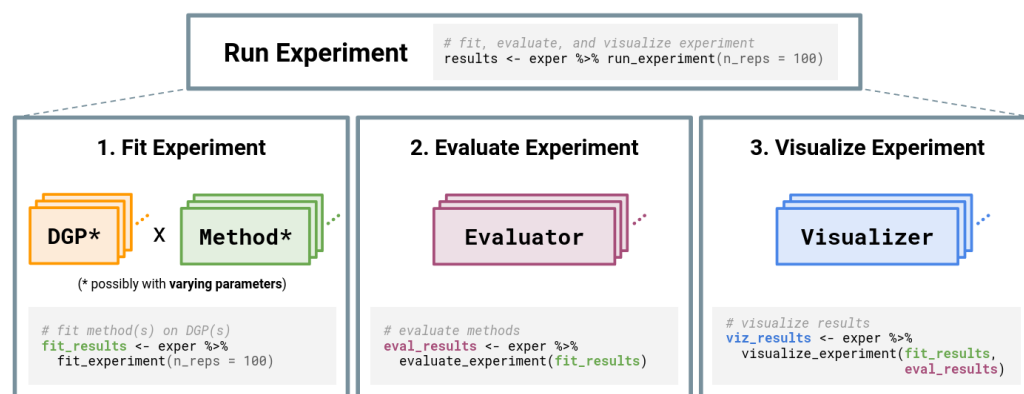


Figure 2: Overview of running a simChef Experiment. The Experiment class handles relationships among the four classes portrayed in Figure 1. Experiments may have multiple DGPs and Methods, which are combined across the Cartesian product of their varying parameters (represented by *). Once computed, each Evaluator and Visualizer takes in the fitted simulation replicates, while Visualizer additionally receives evaluation summaries.

A powerful grammar of data science simulations

Inspired by the tidyverse (Wickham et al., 2019), simChef develops an intuitive grammar for running simulation studies using the aforementioned R6 classes. We provide an illustrative example usage next.

```

library(simChef)

dgp1 <- create_dgp(dgp_fun1, "my_dgp1", sd = 0.5)
dgp2 <- create_dgp(dgp_fun2, "my_dgp2")
method <- create_method(method_fun, "my_method")
eval <- create_evaluator(eval_fun)
viz <- create_visualizer(viz_fun)

exper <- create_experiment(dgp_list = list(dgp1, dgp2)) %>%
  add_method(method) %>%
  add_vary_across(
    list(dgp1, dgp2),
    n = c(1e2, 1e3, 1e4)
  ) %>%
  add_vary_across(
    dgp2,
    sparse = c(FALSE, TRUE)
  ) %>%
  add_vary_across(
    method,
    scalar_valued_param = c(0.1, 1.0, 10.0),
    vector_valued_param = list(c(1, 2, 3), c(4, 5, 6)),
    list_valued_param = list(list(a1=1, a2=2, a3=3),
                              list(b1=3, b2=2, b3=1))
  ) %>%
  add_evaluator(eval) %>%
  add_viz(viz)

future::plan(multicore, workers = 64)
  
```

```
results <- exper %>%
  run_experiment(n_reps = 100, save = TRUE)

new_method <- create_method(new_method_fun, 'my_new_method')

exper <- exper %>%
  add_method(new_method)

results <- exper %>%
  run_experiment(n_reps = 100, use_cached = TRUE)

init_docs(exper)
render_docs(exper)
```

In the example usage, DGP(s), Method(s), Evaluator(s), and Visualizer(s) are first created via `create_*`(). These simulation objects can then be combined into an Experiment using either `create_experiment()` and/or `add_*`().

In an Experiment, DGP(s) and Method(s) can also be varied across one or multiple parameters via `add_vary_across()`. For instance, in the example Experiment, there are two DGP instances, both of which are varied across three values of `n` and one of which is additionally varied across two values of `sparse`. This effectively results in nine distinct configurations for data generation (i.e., 3 variations on `dgp1` + 3x2 variations on `dgp2`). For the single Method in the experiment, we use three values of `scalar_valued_param`, two of `vector_valued_param`, and another two of `list_valued_param`, giving 12 distinct configurations. Hence, there are a total of $9 \times 12 = 108$ DGP-method-parameter combinations in the Experiment.

Thus far, we have simply instantiated an Experiment object (akin to creating a recipe for an experiment). To compute and run the simulation experiment, we next call `run_experiment` with the desired number of replicates. As summarized in Figure 2, running the experiment will (1) *fit* each Method on each DGP (and for each of the varying parameter configurations), (2) *evaluate* the experiment according to the given Evaluator(s), and (3) *visualize* the experiment according to the given Visualizer(s). Furthermore, the number of replicates per combination of DGP, Method, and parameters specified via `add_vary_across` is determined by the `n_reps` argument to `run_experiment`. Because replication happens at the per-combination level, the effective total number of replicates in the Experiment depends on the number of DGPs, methods, and varied parameters. In the given example, there are 108 DGP-method-parameter combinations, each of which is replicated 100 times. To reduce the computational burden, the Experiment class flexibly handles the computation of simulation replicates in parallel using the future package (Bengtsson, 2021). Figure 3 provides a detailed schematic of the `run_experiment` workflow, along with the expected inputs to and outputs from user-defined functions.

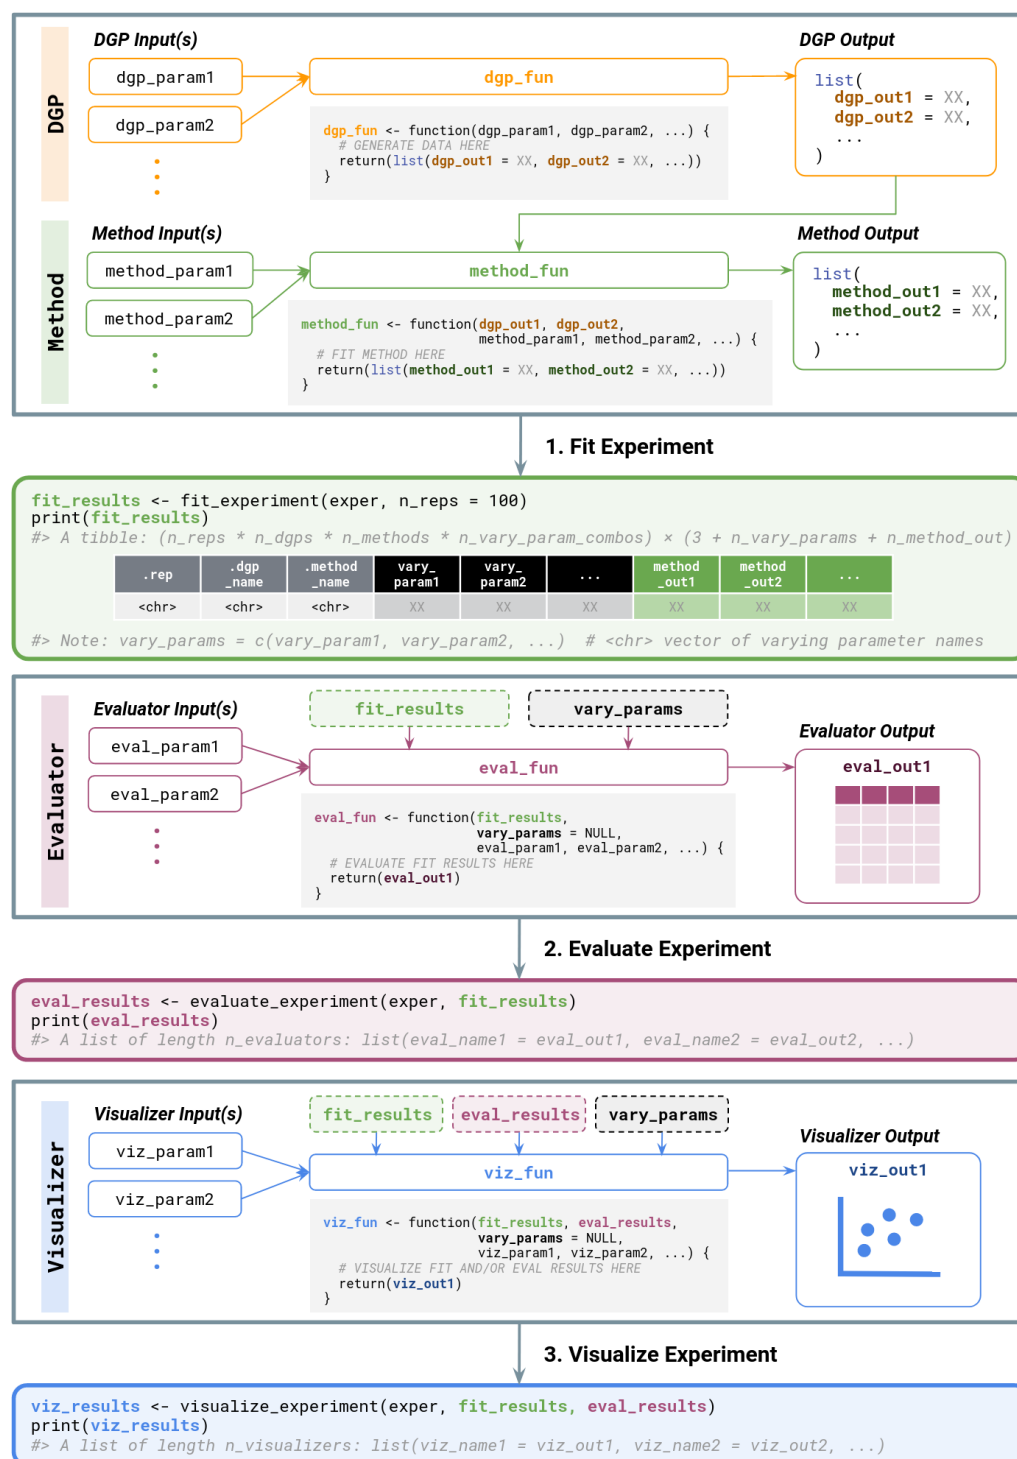


Figure 3: Detailed schematic of the `run_experiment` workflow using `simChef`. Expected inputs to and outputs from user-defined functions are also provided.

Additional Features

In addition to the ease of parallelization, `simChef` enables caching of results to further alleviate the computational burden. Here, users can choose to save the experiment's results to disk by

passing `save = TRUE` to `run_experiment`. Once saved, the user can add new DGP and Method objects to the experiment and compute additional replicates without re-computing existing results via the `use_cached` option. Considering the example above, when we add `new_method` and call `run_experiment` with `use_cached = TRUE`, `simChef` finds that the cached results are missing combinations of `new_method`, existing DGPs, and their associated parameters, giving nine new configurations. Replicates for the new combinations are then appended to the cached results.

`simChef` also provides users with a convenient API to automatically generate an R Markdown document. This documentation gathers the scientific details, summary tables, and visualizations side-by-side with the user's custom source code and parameters for data-generating processes, statistical methods, evaluation metrics, and plots. A call to `init_docs` generates empty markdown files for the user to populate with their overarching simulation objectives and with descriptions of each of the DGP, Method, Evaluator, and Visualizer objects included in the Experiment. Finally, a call to `render_docs` prepares the R Markdown document, either for iterative design and analysis of the simulation or to provide a high-quality overview that can be shared easily. We provide an example of the simulation documentation [here](#). Corresponding R source code is available on [GitHub](#).

Discussion

While `simChef`'s core functionality focuses on computability (C) – encompassing efficient usage of computational resources, ease of user interaction, reproducibility, and documentation – we emphasize the importance of predictability (P) and stability (S) in data science simulations. The principal goal of `simChef` is to provide a tool for data scientists to create simulations that incorporate predictability (through fit to real-world data) and stability (through sufficient exploration of uncertainty) in their simulations. In future work, we intend to provide tools that can be flexibly tailored to a user's particular scientific needs and further these goals through automated predictability and stability summaries and documentation.

Acknowledgements

The authors gratefully acknowledge partial support from (a) the NSF under awards DMS-2209975, 1613002, 1953191, 2015341, and IIS 1741340; and grant 2023505 supporting the Foundations of Data Science Institute (FODSI); (b) the Weill Neurohub; and (c) the Chan Zuckerberg Biohub under an Intercampus Research Award. TMT acknowledges support from the NSF Graduate Research Fellowship Program DGE-2146752.

References

- Bengtsson, H. (2021). A Unifying Framework for Parallel and Distributed Processing in R using Futures. *The R Journal*, 13(2), 208. <https://doi.org/10.32614/RJ-2021-048>
- Chang, W. (2022). *R6: Encapsulated classes with reference semantics*.
- Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Golemund, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M., Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., ... Yutani, H. (2019). Welcome to the Tidyverse. *Journal of Open Source Software*, 4(43), 1686. <https://doi.org/10.21105/joss.01686>
- Yu, B., & Kumbier, K. (2020). Veridical data science. *Proceedings of the National Academy of Sciences*, 117(8), 3920–3929. <https://doi.org/10.1073/pnas.1901326117>